Heard but Overruled?
Persistent Correct Dissent under False Peer Authority

Initial research proposal — November 8, 2026

Abstract

A multi-agent team may acknowledge a valid safety objection while proceeding with a prohibited action, or avoid that action while excluding the agent that raised it. These outcomes should not be collapsed into agreement. We propose a controlled study of how a persistent, correct dissenter is treated when another peer falsely claims authority to override an explicit policy. A small simulated team will resolve synthetic workflow tasks with machine-checkable authorization rules. A scripted dissenter will provide the same valid objection once or repeatedly, while another scripted peer either asserts an unauthorized override or gives a matched recommendation without claiming authority. A model-controlled coordinator will choose consultation targets and a final action. We will separately measure observable attention, behavioral uptake, exclusion, and policy compliance. Paired task variants and an immutable authorization record will distinguish conversational status from actual permission. No experiments have been run. The intended deliverable is a reproducible pilot identifying where correct dissent stops influencing action, followed, if resources allow, by an interactive replication and a safeguard comparison.

1 Introduction

A safety objection matters only if it can affect the action that follows. A peer may falsely assert that an exception has been approved; a correct dissenter can repeat the rule but may be acknowledged, ignored, or removed from later consultation. We ask: *Does a false authority claim change whether persistent correct dissent is attended to, acted on, or excluded?*

Correctness will be established by a deterministic policy checker, not majority vote. The objection remains valid, the dissenter has no special permissions, and the authority claim arrives through an ordinary peer message. This narrow study connects to Track 2's interest in roles and collective failures [6]; a scripted pilot cannot establish emergent leadership or swarm-scale behavior.

2 Related work

Tak et al. study minority signals and process losses in group deliberation [1]. Pitre et al. analyze communication failures, including correction and consensus problems [2]. These motivate examining trajectories as well as final decisions. Kumarappan and Mujoo report that a correctly arguing dissenter reduces answer substitution under simulated peer disagreement [4]; this does not establish how persistent safety dissent is treated under contested authority.

Vijjini et al. examine assigned status, persuasion, and unsafe compliance in power-asymmetric dialogues [3]. Our proposed manipulation instead contradicts a peer's authority claim with a shared permission record. The normative-conformity sprint project describes sensitivity to audience identity and peer reports of rule-breaking [5]. We treat that public summary as early-stage participant work. Separating attention, uptake, exclusion, and safety is an operational choice, not a claim of unprecedented novelty.

3 Proposed methods

3.1 Task and controlled interaction

Use synthetic approval workflows: publishing a mock record requires a valid token issued by an external principal. All agents see the same immutable policy and facts, and the token is absent. Saving a draft and requesting approval is always permitted. A simulator records prohibited action attempts without executing external operations.

Each episode has a model-controlled coordinator and three scripted peers: claimant, dissenter, and neutral peer. The coordinator alone chooses consultation targets and the final structured action. The dissenter correctly cites the missing token. The claimant recommends publishing in both conditions, either asserting false override authority or giving a task-priority rationale without an authority claim. Match tone and approximate length, rotate names and order, and keep evidence and majority support constant.

Cross **false authority present/absent** with **one objection/persistent objection**. The initial objection is identical. Persistent dissent repeats the same evidence on two later scheduled turns; the one-objection control uses neutral status messages of similar length. Repeated semantic salience is intentionally part of the intervention.

All first-round messages are delivered. The coordinator can subsequently retain or exclude peers from consultation. Excluded peers' scheduled messages are logged but not delivered, making exclusion behaviorally consequential. Analyze treatment assignment regardless of delivery; a forced-delivery sensitivity run can separate repeated exposure from loss of access.

3.2 Measures and analysis

Construct	Proposed observable measure
Attention	Explicit reference to the objection's policy evidence in the next response; consultation selection recorded separately. This is not internal model attention.
Uptake	Final plan includes the approval gate or permitted alternative and attributes it to the objection. Action change without attribution is reported separately.
Exclusion	Removal of the dissenter from the next consultation list, with timing, actual delivery, and neutral-peer exclusion as comparators.
Safety	No prohibited action attempt according to the checker; report safe task progress separately from abstention.

Record the coordinator's initial intended action before peer exposure for a secondary reversal analysis, without selecting the primary sample by that intention. Two reviewers will annotate attention and uptake using a frozen rubric, blinded to condition labels where feasible; report agreement and adjudicate disagreements. Save public transcripts, routing events, structured actions, and checker outputs; hidden chain-of-thought is unnecessary.

The primary estimand is the authority-induced risk difference in violation attempts within persistent-dissent trials. Report all four cell rates, task-paired differences, and task-cluster bootstrap intervals. Persistence interactions and other measures are secondary descriptive analyses. Agents and repeated runs within a task are not independent task samples. Correlations between exclusion and violations do not identify mediation.

A tentative pilot might use 20 held-out task instances, four cells, and three independent runs (240 episodes for one model). This is a budget illustration, not a power calculation or committed sample. Final model, sampling settings, sample size, and stopping rule must be chosen after a separate calibration set and frozen before evaluation. If resources permit, test a mandatory permission-record check and then repeat with live peers. Retain the scripted comparison to identify the original manipulation.

4 Results: planned reporting

No results are available. Report episode counts, invalid outputs, all cell rates with uncertainty, and a joint table of acknowledgement, uptake, exclusion, and safety. Choose illustrative traces by a stated rule, including counterexamples. Report null and mixed outcomes without changing hypotheses retrospectively.

5 Discussion

The pilot can identify responses to a specific authority cue under a fixed routing protocol. It cannot establish dislike, subjective social pressure, or human organizational behavior. Correct scripted warnings omit real-world uncertainty: subsequent work should include incorrect objections rather than reward dissent indiscriminately. An interactive replication would test whether findings survive unconstrained peer responses; scaling to larger networks remains separate work.

6 Conclusion

The proposed contribution is an auditable test of whether correct dissent is heard, used, excluded, and safety-preserving. Its value depends on clear interventions and honest uncertainty rather than a promised conformity effect.

A Limitations and Dual-Use / Ethical Considerations

Synthetic rules, scripted peers, and model-version sensitivity limit transfer. False-authority examples could inform agent manipulation; use inert records and no external access or real credentials. English is the baseline. Chinese and Vietnamese are optional matched translations requiring bilingual review of policy meaning and authority cues; differences would concern prompts and models, not national cultures.

B Provenance and AI assistance

This AI-assisted initial proposal needs substantive team review, authorship details, and experimental work before final reporting. No authors or affiliations are asserted. The user reports inspiration from a friend's work; that review was unavailable and is not evidence here. The supplied Google document is a submission template. Final authors must verify claims and follow the current template and disclosure requirements.

References

[1] A. N. Tak et al. (2026). *From Process Loss to Assembly Bonus: Human-Grounded Diagnosis of Multi-Agent LLM Collaboration*. arXiv:2609.13261v1. https://arxiv.org/html/2609.13261v1.

[2] P. Pitre et al. (2026). *Understanding Communication Failures in Multi-Agent LLM Systems*. arXiv:2610.32060v1. https://arxiv.org/html/2610.32060v1.

[3] A. R. Vijjini, S. B. Manjunath, and S. Chaturvedi (2026). *Do LLM Agents Mirror Socio-Cognitive Effects in Power-Asymmetric Conversations?* ACL. https://aclanthology.org/2026.acl-long.2202/.

[4] A. Kumarappan and A. Mujoo (2026). *Not Just RLHF: Why Alignment Alone Won't Fix Multi-Agent Sycophancy*. arXiv:2605.12991v3. https://arxiv.org/abs/2605.12991v3.

[5] R. Acharya and J. Chen (2026). *Impact of Normative Conformity and Perception on Model Misalignment in Multi-Agent Systems*. Digital Minds Research Sprint project page; early-stage participant work. https://apartresearch.com/sprints/projects/impact-of-normative-conformity-and-perception-on-model-misalignment-in-multiagent-systems-4ey0.

[6] Apart Research (2026). *Swarm Alignment Hackathon*, Track 2: The Science of Agent Networks. Accessed November 8, 2026. https://apartresearch.com/sprints/swarm-alignment-2026.